

```

import matplotlib.pyplot as plt
import os
import re
import shutil
import string
import tensorflow as tf
from tensorflow.keras import layers
from tensorflow.keras import losses
print(tf.__version__)

2.15.0

url = "https://ai.stanford.edu/~amaas/data/sentiment/aclImdb_v1.tar.gz"

dataset = tf.keras.utils.get_file("aclImdb_v1", url,
                                  untar=True, cache_dir='.',
                                  cache_subdir='')

dataset_dir = os.path.join(os.path.dirname(dataset), 'aclImdb')

Downloading data from https://ai.stanford.edu/~amaas/data/sentiment/aclImdb_v1.tar.gz
84125825/84125825 [=====] - 6s 0us/step

os.listdir(dataset_dir)

['imdb.vocab', 'test', 'train', 'imdbEr.txt', 'README']

train_dir = os.path.join(dataset_dir, 'train')
os.listdir(train_dir)

['urls_pos.txt',
 'unsupBow.feats',
 'neg',
 'urls_neg.txt',
 'unsup',
 'labeledBow.feats',
 'urls_unsup.txt',
 'pos']

sample_file = os.path.join(train_dir, 'pos/1181_9.txt')
with open(sample_file) as f:
    print(f.read())

Rachel Griffiths writes and directs this award winning short film. A heartwarming story about coping with grief and cherishing the m

```



```

remove_dir = os.path.join(train_dir, 'unsup')
shutil.rmtree(remove_dir)

batch_size = 32
seed = 42

raw_train_ds = tf.keras.utils.text_dataset_from_directory(
    'aclImdb/train',
    batch_size=batch_size,
    validation_split=0.2,
    subset='training',
    seed=seed)

Found 25000 files belonging to 2 classes.
Using 20000 files for training.

for text_batch, label_batch in raw_train_ds.take(1):
    for i in range(3):
        print("Review", text_batch.numpy()[i])
        print("Label", label_batch.numpy()[i])

Review b"Pandemonium" is a horror movie spoof that comes off more stupid than funny. Believe me when I tell you, I love comedies. f
Label 0
Review b"David Mamet is a very interesting and a very un-equal director. His first movie 'House of Games' was the one I liked best,
Label 0
Review b"Great documentary about the lives of NY firefighters during the worst terrorist attack of all time.. That reason alone is v
Label 1

```




McAfee | WebAdvisor

Your download's being scanned.  
We'll let you know if there's an issue.



```
print("1287 ---> ",vectorize_layer.get_vocabulary()[1287])
print(" 313 ---> ",vectorize_layer.get_vocabulary()[313])
print('Vocabulary size: {}'.format(len(vectorize_layer.get_vocabulary())))
```

```
1287 --->  silent
313 --->  night
Vocabulary size: 10000
```

```
train_ds = raw_train_ds.map(vectorize_text)
val_ds = raw_val_ds.map(vectorize_text)
test_ds = raw_test_ds.map(vectorize_text)
```

```
AUTOTUNE = tf.data.AUTOTUNE
train_ds = train_ds.cache().prefetch(buffer_size=AUTOTUNE)
val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)
test_ds = test_ds.cache().prefetch(buffer_size=AUTOTUNE)
```

```
embedding_dim = 16
```

```
model = tf.keras.Sequential([
    layers.Embedding(max_features, embedding_dim),
    layers.Dropout(0.2),
    layers.GlobalAveragePooling1D(),
    layers.Dropout(0.2),
    layers.Dense(1)])
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, None, 16)	160000
dropout (Dropout)	(None, None, 16)	0
global_average_pooling1d (GlobalAveragePooling1D)	(None, 16)	0
dropout_1 (Dropout)	(None, 16)	0
dense (Dense)	(None, 1)	17

=====  
Total params: 160017 (625.07 KB)  
Trainable params: 160017 (625.07 KB)  
Non-trainable params: 0 (0.00 Byte)

```
model.compile(loss=losses.BinaryCrossentropy(from_logits=True),
              optimizer='adam',
              metrics=tf.metrics.BinaryAccuracy(threshold=0.0))
```

```
epochs = 10
```

```
history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=epochs)
```

```
Epoch 1/10
625/625 [=====] - 11s 16ms/step - loss: 0.6623 - binary_accuracy: 0.6956 - val_loss: 0.6113 - val_binary_acc
Epoch 2/10
625/625 [=====] - 6s 9ms/step - loss: 0.5447 - binary_accuracy: 0.8048 - val_loss: 0.4949 - val_binary_acc
Epoch 3/10
625/625 [=====] - 4s 7ms/step - loss: 0.4410 - binary_accuracy: 0.8478 - val_loss: 0.4179 - val_binary_acc
Epoch 4/10
625/625 [=====] - 5s 8ms/step - loss: 0.3767 - binary_accuracy: 0.8670 - val_loss: 0.3725 - val_binary_acc
Epoch 5/10
625/625 [=====] - 4s 6ms/step - loss: 0.3342 - binary_accuracy: 0.8803 - val_loss: 0.3444 - val_binary_acc
Epoch 6/10
625/625 [=====] - 4s 6ms/step - loss: 0.3033 - binary_accuracy: 0.8905 - val_loss: 0.3251 - val_binary_acc
Epoch 7/10
625/625 [=====] - 8s 13ms/step - loss: 0.2808 - binary_accuracy: 0.8986 - val_loss: 0.3124 - val_binary_acc
Epoch 8/10
625/625 [=====] - 8s 12ms/step - loss: 0.2606 - binary_accuracy: 0.9086 - val_loss: 0.3000 - val_binary_acc
Epoch 9/10
625/625 [=====] - 5s 8ms/step - loss: 0.2453 - binary_accuracy: 0.9186 - val_loss: 0.2875 - val_binary_acc
Epoch 10/10
625/625 [=====] - 4s 6ms/step - loss: 0.2305 - binary_accuracy: 0.9286 - val_loss: 0.2750 - val_binary_acc
```



McAfee | WebAdvisor

Your download's being scanned.  
We'll let you know if there's an issue.

```

loss, accuracy = model.evaluate(test_ds)
print("Loss: ", loss)
print("Accuracy: ", accuracy)

782/782 [=====] - 4s 5ms/step - loss: 0.3103 - binary_accuracy: 0.8730
Loss: 0.3102893829345703
Accuracy: 0.8729599714279175

```

```

history_dict = history.history
history_dict.keys()

dict_keys(['loss', 'binary_accuracy', 'val_loss', 'val_binary_accuracy'])

```

```


acc = history_dict['binary_accuracy']
val_acc = history_dict['val_binary_accuracy']
loss = history_dict['loss']
val_loss = history_dict['val_loss']

epochs = range(1, len(acc) + 1)

# "bo" is for "blue dot"
plt.plot(epochs, loss, 'bo', label='Training loss')
# b is for "solid blue line"
plt.plot(epochs, val_loss, 'b', label='Validation loss')
plt.title('Training and validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.show()

```



```

plt.plot(epochs, acc, 'bo', label='Training acc')
plt.plot(epochs, val_acc, 'b', label='Validation acc')
plt.title('Training and validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend(loc='lower right')

plt.show()


```



McAfee | WebAdvisor



Your download's being scanned.  
We'll let you know if there's an issue.



```
export_model = tf.keras.Sequential([
    vectorize_layer,
    model,
    layers.Activation('sigmoid')
])
export_model.compile(
    loss=losses.BinaryCrossentropy(from_logits=False), optimizer="adam", metrics=['accuracy']
)
# Test it with `raw_test_ds`, which yields raw strings
loss, accuracy = export_model.evaluate(raw_test_ds)
print(accuracy)

782/782 [=====] - 5s 6ms/step - loss: 0.3103 - accuracy: 0.8730
0.8729599714279175
```

```
examples = [
    "The movie was great!",
    "The movie was okay.",
    "The movie was terrible..."
]
export_model.predict(examples)

1/1 [=====] - 0s 218ms/step
array([[0.60859257],
       [0.4288      ],
       [0.34703887]], dtype=float32)
```



McAfee | WebAdvisor



Your download's being scanned.  
We'll let you know if there's an issue.